

# Hardware Implementation of Neural Network Inference Using SystemVerilog - Phase 2: Fixed-Point Neural Computation

## 1. Introduction

---

This phase continues our previous work with designing and implementing a neural network and converting those trained parameters to fixed-point for further inference in SystemVerilog. To accomplish this, a shallow neural network was implemented from scratch in Python. The trained parameters were extracted and converted to fixed-point format for use in SV. The dataset used for this project is called the Wisconsin Breast Cancer Dataset (WBCD), which can be imported from the ``sklearn.datasets`` package.

This phase consists of three sections.

Section two focuses on Data Selection, with a primary focus on analyzing the dataset, deriving useful information through statistical analysis and visualization, and selecting the best feature pair based on computed metrics and measurements.

Section three focuses on implementing the Neural Network Class from scratch and training it using the two selected, scaled features, and retrieving the trained parameters for further quantization investigations.

Finally, in section four, the importance of fixed-point representation in SystemVerilog will be discussed, and parameters will be converted using an appropriate quantization scale.

## 2. Data Analysis

---

This section starts by displaying a sample of data, providing a descriptive summary, and visualizing data. We will then proceed with selecting the best features and performing feature engineering.

### 2.1 Dataset

---

The first step involves loading in and analyzing the data using ``Pandas``, ``NumPy``, ``Matplotlib``, and ``Seaborn``. The original data has about 30 features and 569 records (examples).

A sample of the data is shown below:

| mean radius | mean texture | worst smoothness | worst texture | worst perimeter |
|-------------|--------------|------------------|---------------|-----------------|
| 17.99       | 10.38        | 0.1622           | 17.33         | 184.60          |
| 20.57       | 17.77        | 0.1238           | 23.41         | 158.80          |

|       |       |        |       |        |
|-------|-------|--------|-------|--------|
| 19.69 | 21.25 | 0.1444 | 25.53 | 152.50 |
|-------|-------|--------|-------|--------|

Table 1. A sample of Data

## 2.2 Descriptive Summary

The number of null rows has been checked using `data.info()`. It is worth noting that no NaN values were found. Also, a descriptive summary showed the distribution of the features using `data.describe()` method.

## 2.3 Correlation Analysis

Since 30 features are too many for a small network with the architecture  $2 \rightarrow 4 \rightarrow 2$ , only the two most correlated features have been utilized as the input. The objective is to identify features with low pairwise correlation and high Pearson correlation with the target. A heatmap has been plotted to show the correlation between features and the target. (Fig. 1)




Figure 1. Correlation Heatmap

## 2.4 Feature Selection

---

Since choosing the two most correlated features based on only a Heatmap would be a difficult choice, quantitative metrics such as **Pearson correlation** (from the `SciPy` package) and the **feature-pair scoring** method have also been utilized to make sure the selected features are scientifically reasonable.

The feature-pair scoring has been computed using the following formula:

$$\text{score} = |\text{corr}(A, y)| |\text{corr}(B, y)| (1 - |\text{corr}(A, B)|)$$

Where:

- $\text{corr}(A, y)$ : correlation between feature A and target
- $\text{corr}(B, y)$ : correlation between feature B and target
- $\text{corr}(A, B)$ : correlation between the two features

Top 3 features with the lowest multicollinearity and highest Pearson correlation:

| feature_1      | feature_2        | target_corr_1 | target_corr_2 | feature_corr | score    |
|----------------|------------------|---------------|---------------|--------------|----------|
| mean radius    | worst smoothness | 0.730029      | 0.421465      | 0.119616     | 0.270878 |
| mean perimeter | worst smoothness | 0.742636      | 0.421465      | 0.150549     | 0.265874 |
| mean area      | worst smoothness | 0.708984      | 0.421465      | 0.123523     | 0.261902 |

Table 2. Correlation Map between Features and Target

“*worst smoothness*” appears alongside other features in all three sets. The pair of features in the first row has the best performance, since the score is higher and the feature\_corr is lower. Higher scores indicate feature pairs that simultaneously reveal strong correlations with the target and relatively low redundancy with each other. It is worth mentioning that for further analysis, training the network with all three pairs might end up in a more accurate result. This experiment is especially worth it in selecting the best pair of features. As an example, *mean perimeter* in the second pair has a higher correlation than the *mean radius* in the first pair. But this matter should also take into consideration that the higher the correlation of the features, the lower the final score. Therefore, selecting the best pair of features is actually a trade-off that should be handled in further analysis.

## 2.5 Outlier Analysis

---

Further visualization techniques, such as box plots, revealed that the selected features include outliers, which are data points that fall outside a defined range. Neural Networks can be sensitive to outliers, especially when training data are limited. Therefore, further analysis has been conducted, and various functions have been implemented to detect them. Since several outliers appear to belong to clinically relevant classes rather than measurement errors, they were retained to avoid discarding potentially informative observations.

A box plot is shown below (Fig. 2):




Figure 2. Boxplot for Displaying Data Distribution

## 2.6 Feature Scaling

To address the differences in variable ranges as shown in Table 1, a Scaling method, `MinMaxScaler(feature_range=(-1, 1))`, has been applied to the features after splitting them into train-test sets, ensuring they fall within the range of  $[-1, 1]$ .

| Info  | worst smoothness | mean radius |
|-------|------------------|-------------|
| count | 569              | 569         |
| mean  | 0.132369         | 14.127292   |
| std   | 0.022832         | 3.524049    |
| min   | 0.071170         | 6.981000    |
| 25%   | 0.116600         | 11.700000   |
| 50%   | 0.131300         | 13.370000   |
| 75%   | 0.146000         | 15.780000   |
| max   | 0.222600         | 28.110000   |

Table 3. Descriptive info of unsclaed features

### 3. Model Implementation

---

This section continues with implementing a neural network from scratch using Python and NumPy. We will go through data splitting, scaling methods, tensor conversion, and network architecture.

#### 3.1 Data Splitting

---

The process of training a model begins by splitting the clean data into two sets: training and testing. If there is sufficient data, it is also beneficial to create a validation set. The analysis revealed that a class imbalance in the target variable exists. Therefore, the `stratify` parameter was set to `y` (the target). It is worth noting that 10% of the entire dataset was reserved for the test set, while the remaining 90% was allocated for training.

#### 3.2 Data Preprocessing

---

The network needs scaled, preprocessed training and testing sets before being fed. `X_train` should be fit and transformed, while `X_test` should only be transformed. A data transpose is also required for the `X` values for further matrix multiplication computations. Target values should be one-hot encoded.

#### 3.3 Neural Network Architecture

---

The architecture of neural networks includes several key components. It starts with the input layer, which receives the raw data. This is followed by one or more hidden layers, where neurons process the input using activation functions like sigmoid, ReLU, or tanh to introduce non-linearity. Finally, the output layer generates the network's predictions or classifications based on the processed information. Weights connect the neurons and are adjusted during training to improve the network's performance.

#### 3.4 Activation Functions

---

The two activation functions used in this project are mainly 'ReLU' (also its derivative for backpropagation) and 'Softmax'.

### 3.5 Methods

---

Each of the steps in a neural network, from initializing the parameters to updating them, was written as a single method, each responsible for a certain task. This is known as the Single Responsibility Principle (or SRP). Please note that the network is implemented from scratch (using NumPy and pure Python), since it is easier to retrieve and store the weights/biases for smaller datasets.

Methods of the Neural Network class include: `'initialize_parameters'`, `'forward_propagation'`, `'compute_cost'`, `'backward_propagation'`, `'initialize_adam'`, `'update_parameters'`

### 3.6 Class Neural Network

---


The network is written using a class and has methods that are in a separate file called `'methods.py'`. The general purpose of this class is to train the network and thus get the best set of parameters (i.e., updated parameters retrieved from the Gradient Descent). The hyperparameters like `'learning_rate'`, `'decay_rate'`, `'penalty'`, and `'epochs'` were selected using `'GridSearchCV'`.

### 3.7 Result

---

The training resulted in 93.75% accuracy on the training set and 92.98% on the testing set. The model correctly classified 53 out of 57 test samples. Indices with the wrong predictions are array([ 2, 5, 40, 56]). Further analysis is required to determine whether the misclassified samples correspond to outliers or overlapping class regions.

A visual representation of the training's cost function has been provided below (Fig. 3):



### 3.8 Evaluation Metrics

Below, some useful evaluation metrics have been provided that show the overall network performance (Fig. 4). These include: *classification report & confusion matrix*.

*Classification Report:*

| Info         | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| malignant    | 0.90      | 0.90   | 0.90     | 21      |
| benign       | 0.94      | 0.94   | 0.94     | 36      |
| accuracy     |           |        | 0.93     | 57      |
| macro avg    | 0.92      | 0.92   | 0.92     | 57      |
| weighted avg | 0.93      | 0.93   | 0.93     | 57      |

Table 4. Classification Report

*Confusion Matrix:*

[[19 2]  
[ 2 34]]




Figure 4. Confusion Matrix

The network achieved higher precision and recall for benign samples than for malignant samples.

The first 11 test data points along with their classes are shown below:

| Sample # | Prediction | Ground Truth |
|----------|------------|--------------|
| 0        | 0          | 0            |
| 1        | 1          | 1            |
| 2        | 1          | 0            |
| 3        | 0          | 0            |
| 4        | 0          | 0            |
| 5        | 0          | 1            |
| 6        | 0          | 0            |
| 7        | 1          | 1            |
| 8        | 1          | 1            |
| 9        | 1          | 1            |
| 10       | 1          | 1            |

Table 5. Eleven Predictions along with their True Labels

### 3.9 Parameters

---

Since the main focus of this project is on designing a hardware-based simulated network in SystemVerilog, for the sake of simplicity, the weights/biases are converted to **fixed-point** rather than **floating-point**. Note that all parameters, weighted sum, and the activations are all required to return the maximum among them for input quantization. The first four tables show the parameters in floating-point format.

- a. Floating-point parameters:



**First-layer weight matrix ( $W_1$ )**

| Hidden Neuron | worst smoothness | mean radius |
|---------------|------------------|-------------|
| h_1           | 0.05276146       | 0.04545019  |
| h_2           | 0.87860578       | 1.4938967   |
| h_3           | 0.04193307       | 0.03875046  |
| h_4           | 0.84913852       | 1.46507294  |

**First-layer bias vector ( $b_1$ )**

| Hidden Neuron | Bias        |
|---------------|-------------|
| h_1           | -0.0500322  |
| h_2           | 0.93571265  |
| h_3           | -0.05000003 |
| h_4           | 0.90332097  |

**Second-layer weight matrix ( $W_2$ )**

| Output Neuron | h_1        | h_2         | h_3         | h_4         |
|---------------|------------|-------------|-------------|-------------|
| O_1           | -0.0343271 | 1.42355861  | -0.05331349 | 1.46707236  |
| O_2           | 0.01890401 | -1.42651996 | 0.0436103   | -1.45686225 |

**Second-layer bias vector ( $b_2$ )**

| Output Neuron | Bias        |
|---------------|-------------|
| O_1           | -1.20539659 |
| O_2           | 1.20539659  |

b. Fixed-point parameters:

**First-layer weight matrix ( $W_1$ )**

| Hidden Neuron | worst smoothness | mean radius |
|---------------|------------------|-------------|
| h_1           | 216              | 186         |
| h_2           | 3599             | 6119        |
| h_3           | 172              | 159         |
| h_4           | 3478             | 6001        |

**First-layer bias vector ( $b_1$ )**

| Hidden Neuron | Bias |
|---------------|------|
| h_1           | -205 |
| h_2           | 3833 |
| h_3           | -205 |
| h_4           | 3700 |

**Second-layer weight matrix ( $W_2$ )**

| Output Neuron | h_1  | h_2   | h_3  | h_4   |
|---------------|------|-------|------|-------|
| O_1           | -141 | 5831  | -218 | 6009  |
| O_2           | 77   | -5843 | 179  | -5967 |

**Second-layer bias vector ( $b_2$ )**

| Output Neuron | Bias  |
|---------------|-------|
| O_1           | -4937 |
| O_2           | 4937  |

#### 4. Fixed-Point Parameters

---

To make predictions and run inferences in HDL / SystemVerilog, the parameters needed to be converted from floating-point to fixed-point and stored in n-bit signed signals.

A **16-bit signed fixed-point** representation was selected because:

- Easy to implement in hardware
- Has a small memory footprint
- Provides sufficient precision for many machine-learning applications

To determine the fixed-point format, the following measurements should be computed in order:

- The maximum observed magnitude among parameters,  $X_{train}$ ,  $X_{test}$ , weighted sums, and activations denoted as  $M$
- The required number of integer bits, which is denoted as  $m$
- The required number of fractional bits, which is denoted as  $n$
- $Q_{m.n}$

In this work:

- $M = \max(|Z_2|) \approx 5.131$
- $2^m > M \Rightarrow m = \lceil \log_2(M) \rceil = \lceil \log_2(5.131) \rceil = 3$
- In 16 bits:  $1 + m + n = 16$
- Therefore:  $n = 12$

The computation yields **Q3.12**, which supports the range  $[-8, 7.999755859375)$ , while still preserving good precision using 12 fractional bits.

The corresponding scaling factor is:

$$Scale = 2^{12} = 4096$$

The floating-point values have been converted to fixed-point using the following formula:

- Quantization scale:  $2^{12} = 4096$
- $X_{fixed} = round(x * 4096)$

It is worth mentioning that quantized inference using Q3.12 achieved the same testing accuracy (92.98%) as the floating-point implementation, confirming that Q3.12 provides sufficient numerical precision and dynamic range for the neural-network inference task.

## Summary

---

In this phase, the implementation of a shallow neural network for the Wisconsin Breast Cancer Dataset was covered with a focus on converting trained parameters to fixed-point format for SystemVerilog. Data analysis and feature selection were conducted using correlation metrics, which resulted in the selection of "mean radius" and "worst smoothness" as key features with the lowest multicollinearity and highest correlation to the target.

## References

---

A. Ng, "Neural Networks and Deep Learning," [DeepLearning.AI](https://www.coursera.org/learn/neural-networks-deep-learning), Coursera. [Online]. Available: <https://www.coursera.org/learn/neural-networks-deep-learning>. [Accessed: Jun. 20, 2026].

By Kiana Jafari - Computer Engineering Student